UAE HACKATHON 2026 · TRACK 2 PIONEERS · CHALLENGE #27

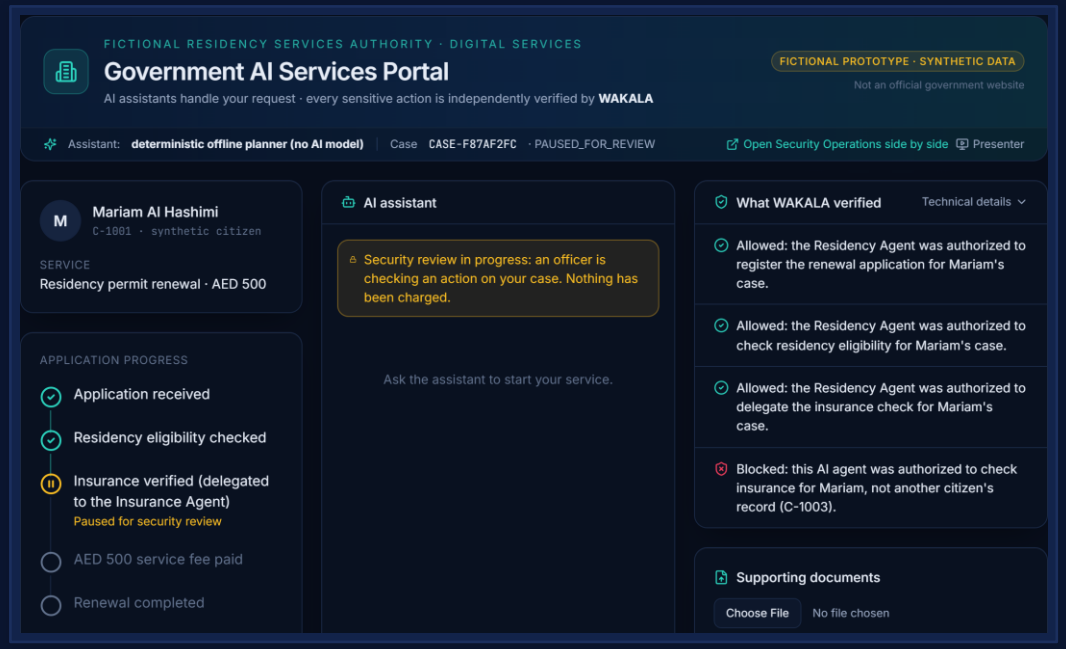
WAKALA وكالة

Autonomous AI. Accountable Actions.

An independent authorization gateway that makes every government AI agent identifiable, bound to one citizen's case, stoppable in seconds and safely recoverable.

Secure Agentic AI for Government Services

WORKING PROTOTYPE · SYNTHETIC DATA · NOT PRODUCTION



Real prototype screen: a delegated AI agent asked for another citizen's record — WAKALA blocked it before it ran.

THE CHALLENGE

When an AI agent acts for government, who is accountable?

Challenge #27 (Cyber Security Council) frames a national shift of **half of government services and operations to agentic AI within two years**. Agents plan, call tools and delegate to other agents on their own.



Who is acting?

Agents and sub-agents run on broad roles and reusable credentials — not as accountable, identifiable actors.

Root causes R1, R2



What may it do?

Delegated authority is not translated into precise limits: which citizen, which tool, how much, until when.

Root cause R3



How do we stop it safely?

Detection, intervention and recovery sit in separate teams. Stopping one bad agent can stop the whole service.

Root cause R4



Result: an erroneous, manipulated or compromised agent can expose citizen data, move money or interrupt a public service — and no one can prove afterwards who authorized it.

HOW IT WORKS

An independent gateway between every agent and every service



Behaviour Watch

Rule-based signals; graduated response

Officer intervention

Suspend one mandate, reassign, resume

Tamper-evident audit

Hash-chained record of every decision

Document Shield

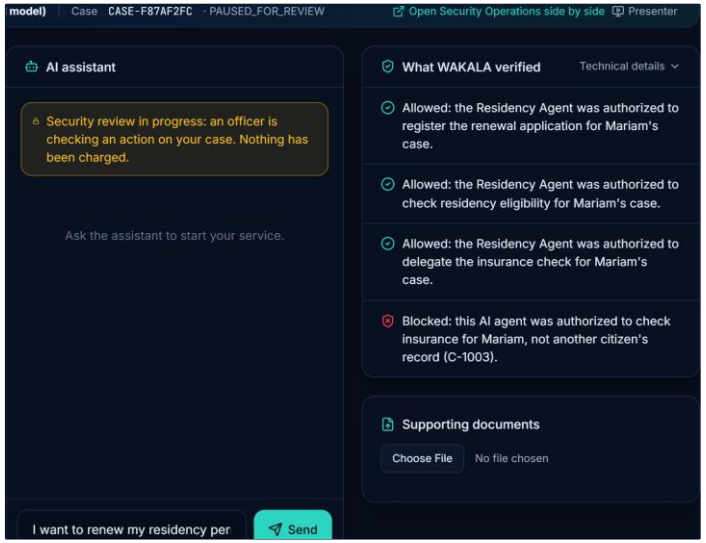
Quarantines hidden prompt-injection PDFs

Citizen AI receipt

What each agent did, in plain language

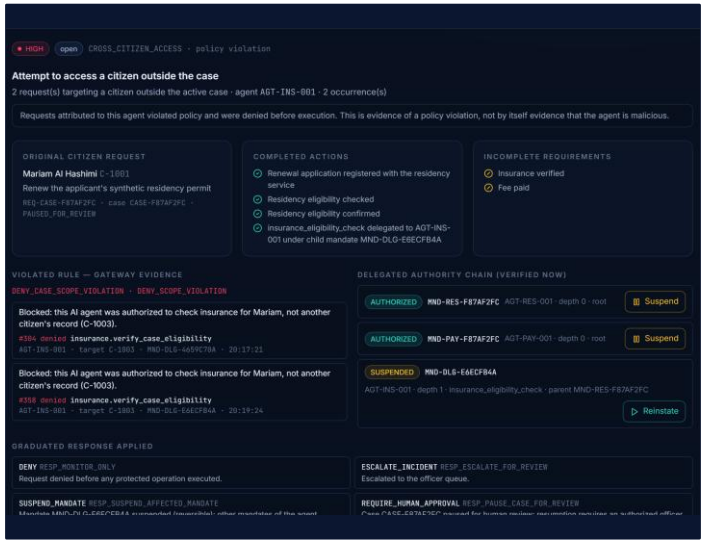
DEMONSTRATION

Live prototype: blocked, contained, recovered



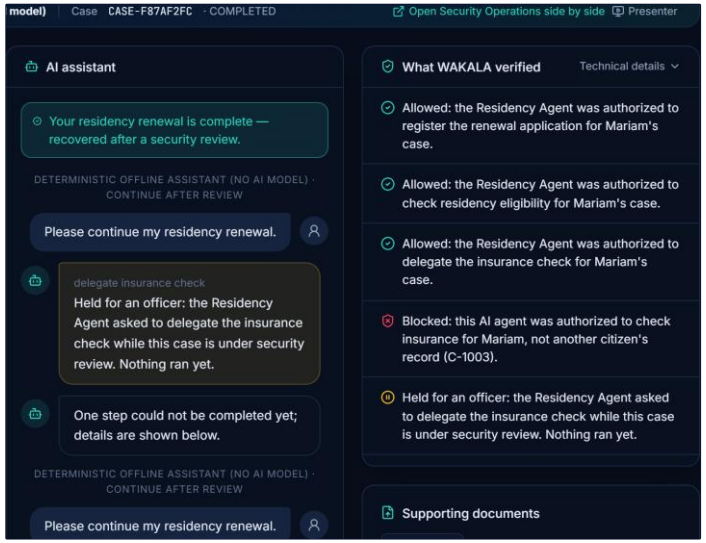
1 Blocked before execution

The delegated Insurance Agent asks for citizen C-1003's record. Mariam's case is paused; nothing charged.



2 Contained precisely

Security Operations: violated rule, verified chain — only the one child mandate is suspended.



3 Service recovered

Officer reassigns to a standby agent; the renewal completes “recovered after a security review”.

Same run: an unrelated citizen's case completed · audit chain verified (386 events) · plain-language AI-activity receipt for Mariam

EVIDENCE

Measured, not claimed

0 / 720

unauthorized operations executed

Full WAKALA · intentionally weak role-only baseline executed 624 / 720

0 / 480

legitimate requests denied

Same 1,200-request labelled synthetic dataset, fixed seed

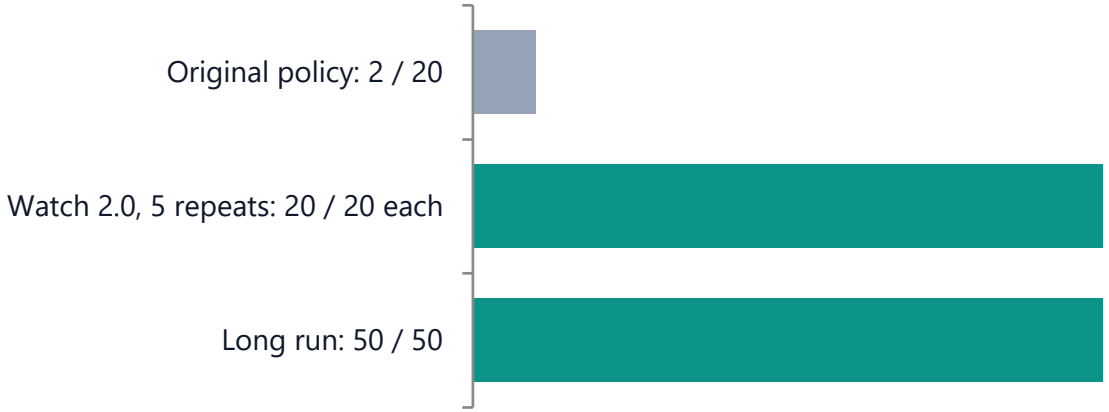
0 / 8

Attack Lab attacks executed

Baseline executed 7 / 8; one (delegation escalation) denied by both

Synthetic prototype benchmark — not production government statistics. Sources: evidence/phase5, phase6, phase7; docs/FINAL_TECHNICAL_READINESS.md.

Safe recovery, shared environment (back-to-back incidents)



Root cause found and fixed: agent-wide counters attached incidents to unrelated cases. Every attack was still denied in every run.

≈19 ms

median full gateway authorization (in-memory; ≈55 ms with persisted SQLite) on one laptop

493 tests

backend tests pass, incl. 34 security invariants; 12 / 12 demo steps verified end-to-end

BUSINESS MODEL (PROPOSED)

Who pays, and why they would

Customers

Government entities deploying agents
Accountable for service outcomes, lawful processing and continuity

Central digital-government & cyber authorities
Need common, cross-entity assurance and evidence

Agent-platform suppliers & integrators
Must show equivalent protection to win government work

Value

Prove every agent action was authorized
Case-bound mandates + hash-chained audit

Stop one agent, not the service
Selective suspension, officer reassignment, recovery

Platform-neutral control
Sits at the service API, not inside one vendor's model

Proposed licensing

Entity licence (annual)
Tiered by protected services and registered agents

Supplier conformance kit
Attack Lab scenarios + evidence pack for platform vendors

Integration & assurance services
Onboarding, mandate templates, officer playbooks

Channels: national shared-service procurement and system-integrator partners. Pricing is a proposal to be validated in the pilot — no revenue or market figures are claimed.

90-day pilot: observe first, then enforce, then decide

Days 1–30 Observe

- Select one low-risk service, its owner and an accountable sponsor
- Agent inventory and mandate templates with the service owner
- Deploy in the entity's sandbox; replace dev keys with entity key management
- Shadow Mode on copied traffic — nothing is blocked

Shadow report: agreement, would-block and false-denial candidates

Days 31–60 Enforce one case type

- Gateway enforcing on one case type; officer approval for payments
- Tune Behaviour Watch thresholds on pilot data
- Train case officers on suspend / reassign / resume
- Independent security test using Attack Lab scenarios

Zero unauthorized actions executed; false denials within agreed limit

Days 61–90 Recover and decide

- Delegation and selective-revocation drills
- Timed recovery drills: incident → safe resume
- Citizen receipt usability test; audit export to the entity's SOC
- Evaluation report to sponsor

Go / no-go decision on measured results

Prerequisites from the pilot entity: a sandbox with test data, a service owner, an officer team, and the entity identity provider and key management.

HONEST STATUS AND OUR ASK

Ready to pilot — honest about what is not built yet

Working today

- Signed identities, case-bound mandates, verified delegation
- Behaviour Watch, officer intervention, safe recovery
- Attack Lab, Document Shield, Shadow Mode, citizen receipt
- Tamper-evident audit; 493 passing tests

Not yet

- No live AI model verified — demo uses a labelled offline planner
- No government integration; synthetic data only
- Single process, SQLite, development keys
- Policy drafts are signed but not enforced

Our ask

- One government entity and service owner as pilot partner
- A sandbox with test data and identity integration
- Joint security review aligned with national cyber guidance
- 90 days to a measured go / no-go

WAKALA وكالة Autonomous AI. Accountable Actions.